

CS170 Lab 2 - Multiprogramming with KOS

 sites.cs.ucsb.edu/~rich/class/cs170/labs/kos_mp/index.html

CS170 KOS


- CS170 Lab Assignment 2 -- Processes and Multiprogramming
- [Rich Wolski](#) and [James Plank](#)
- URL: http://www.cs.ucsb.edu/~rich/class/cs170/labs/kos_mp
- Directory: [/cs/faculty/rich/public_html/class/cs170/labs/kos_mp](#)
- [Culinary advice](#) from a previous chef
- An example [Makefile](#) that will work for submitting your solution
- You must include [simulator_lab2.h](#) to pick up the right set of simulator defines for this lab.
- A C function for printing the initial stack: [printstack.c](#)
- **Due: Wednesday Nov 16th, 2022 at 11:59 PM**

Lab 2 -- Simple t

In this lab, you will add the word "add" in the p may find yourself tuning way. Each piece you c on *Lab 1* parts to make Specifically, you will d to eight user processes can use to perform tim

Also, you will be able **malloc()** and the stan

When you are done, y



Notice the use of working well, you nt works in this self spending time eality.

and you will allow up a timer that you

that make use of

lls:

- **ioctl()** (in limited form)
- **fstat()** (in limited form)
- **getpagesize()**
- **sbrk()**
- **execve()**
- **getpid()**
- **fork()**
- **getppid()**
- **wait()**
- **getdtablesize()**
- **close()** (in limited form)

Moreover, you'll be able to execute a shell and use it to run user programs (although file redirection and pipes won't work).

Changes to the simulator

- For this lab, you need to include the version of **simulator_lab2.h** from this directory. Note that this is different from the **simulator.h** in lab 2. You will need to link your code with **/cs/faculty/rich/cs170/lib/main_lab2.o** and **/cs/faculty/rich/cs170/lib/libsim.a**. There is a Makefile in this directory.

- For this lab, there are two special simulator variables:

```
extern int User_base;  
extern int User_limit;
```

These define what **KOS** addresses the user sees as its memory. When a user job is executing, its address space is mapped to **User_Limit** by the simulator. The user can use up to **User_Limit** bytes of memory. The user process cannot access any memory beyond **User_Limit** and **(User_Base + User_Limit)**. Otherwise, the user process will crash.

- User_Base** and **User_Limit** define the user's memory space. The user program() can use up to **User_Base + User_Limit** bytes of memory. It returns -1 on failure. You can use this as a failure signal.
- There is now a **start_timer(int ticks)**, which tells the user process to start a timer. The default is no timer.

Your Job in this

Basically, we want to do four things in this lab:

- Allow the user to compile and run programs that use the standard I/O library and **malloc()**.
- Implement the parts of KOS that enable a simple shell to work: **fork()**, **execve()** and **wait()**.
- Implement process id's.
- Make time-slicing work.

I will describe each of these in turn.

Allow the user to compile and run programs that use the standard I/O library and malloc()

As before, set up **KOS** so that it loads the program **a.out** as its user process when it starts up. What we first want to do is allow these **a.out** programs to use the standard I/O library and **malloc()**. To do this, we need to do some busy work and implement some system calls that normally we would not care about. The first few are: **getpagesize()**, **getdtablesize()**, and a simple **close()**. **Getpagesize()** should return the value of the **PageSize** variable in **simulator.h**. **Getdtablesize()** should return 64. **Close()** should be implemented so that it returns an error (that's not really how to implement **close()**, but it will suffice for this lab).


We also have to implement one case of **ioctl()** and one case of **fstat()**. Read the man page on **ioctl()**. I don't expect you to understand much about **ioctl()** except for its syntax. You are going to need to implement **ioctl()** when the first argument is 1 and the second argument is **JOS_TCGETP**. Your job is to fill in the third argument which is a pointer to a **(struct JOSTermios)**. You do this by calling **ioctl_console_fill(char *addr)**, where **addr** is the KOS address of the third argument. Then return zero. This is probably confusing, but it must be done.

Fstat() is called by the user process to get file descriptors 0, 1 and 2. You should implement **stat_buf_fill(char *addr)** that the user passed to **fstat()** that the file descriptor should allow one and two, try 256.

Last, you must implement the stack and the heap. You should allow to increase page size.

When you have implemented the stack and the heap, you should allow to use the stack. You should implement the programs in

[/cs/faculty/rich/cs170](#)



should buffer I/O on file descriptors 0, 1 and 2. You should implement **stat_buf_fill(char *addr)** that the user passed to **fstat()** that the file descriptor should allow one and two, try 256.

detail. As before, **close()** should not be

ms should be **close()**. Look at the

, and try ones like [hw2.c](#) that use **printf()**. You can write, compile and run your own programs too -- just follow the compilation steps in [Makefile.xcomp](#) and the instructions from [lab 1](#).

Implement the parts of KOS that enable a simple shell to work: **fork()**, **execve()** and **wait()**

This means you must do multiprocessing. In the last lab, you allowed one user program to execute and gave it all of memory. This time, you should divide memory into 8 equal parts and when you create a new process, it will use one of those parts. Context switching will now involve saving/restoring the registers, **User_base** and **User_limit**.

You'll need to modify the **exit()** system call so that instead of halting **KOS**, it simply terminates a process. **Fork()** and **execve()** are rather straightforward. You should ignore the third argument to **execve()** -- we won't have any environment variables. You also must implement **wait()** and process id's, which should work as in Unix. In particular, processes have parents; otherwise they are orphans. If a process exits and its parent hasn't called **wait()** it becomes a zombie -- it releases its memory but maintains a PCB until its parent either dies or calls **wait()**. Orphans must do the correct thing when they die.

To test this, run

/cs/faculty/rich/cs170/test_execs/ksh

in your version of kos.

WARNING: This is a simple shell with few features. You do not need to use it and can instead, write your own test codes.

The *ksh* binary is capable of running hw from ksh thus:

```
./kos -a /cs/faculty/rich/cs170/test_execs/ksh
```

and you should see

```
KOS booting... done.
Probing console... done
ksh:
```

now run *hw*

```
ksh:/cs/faculty/rich/cs170/test_execs/ksh
Hello world
the write statement
ksh:
```



can run hw from

The *ksh* binary forked and executed the *hw* binary and called **wait()** to wait for it to complete before printing the *ksh:* prompt again.

This binary has few features and many bugs. You can use it to test fork, exec, and wait or you can write your own test codes. If it doesn't work for you, then you must write your own test codes (which is a good idea anyway). The TAs will not be able to answer questions about *ksh* -- either it works for you and you trust it or it doesn't and/or you don't in which case write your own testers.

Implement process id's

Finally, you need to implement **getpid()** and **getppid()** to work with your process id's. This is straightforward. I have no process with process id 0. Orphans return pid 0 as their parent.

- scheduler.h
- syscall.c
- syscall.h

The submission acceptance program will be looking for these files. You should design your solution so that it is contained in files with these names.

This requirement has a few ramifications. First, this list is the file list -- not a subset of the file list. Thus, you can't include extra files and then use "#include" to include those files in these file names. The submission program will use these files and a prepared Makefile to build your program. It will not be prepared for hidden dependencies introduced by extra included files.

This requirement also means that you simply can't concatenate a bunch of files that contain your solution to create one of more of these files. The concatenation is likely to cause double definitions as header files are included multiple times which will cause your compile to fail.

Header files and libraries have been used in the past. If you have been using in class (like libfdr, kt.h, etc.) then you should use the class makefile or use the example included here

